

AIRTOS: Evidence-Constrained Admission, Resource Governance, and Recovery for Heterogeneous Edge AI

Anonymous Authors

Abstract—Heterogeneous edge-AI execution spans CPU, vector, accelerator, and DMA resources while sharing bounded memory and non-coherent buffers. A compiler-generated plan can therefore be structurally valid yet inapplicable to the current target, evidence policy, provider state, memory state, arrival model, or recovery budget. This paper presents AIRTOS, an evidence-constrained runtime governance layer that treats a compiler plan as a first-class admission object. AIRTOS evaluates a joint predicate over package parsing, binding, applicability domain, evidence coverage, provider availability, memory, coherency, schedulability, and recoverability. It combines generation-checked arena allocation with conservative multi-resource schedule simulation in one atomic admission transaction, dispatches dependency-ready segments through stable per-resource earliest-deadline-first queues, executes plan-declared cache ownership transitions, rejects stale completions by device epoch and dispatch cookie, and re-applies the complete gate before fallback. We evaluate these mechanisms in four claim-oriented experimental cores and on one CanMV-K230-LP4 V3.0 board running RT-Smart. The frozen corpus produced no observed safety failure in 3,900 admission cases, 400,000 concurrent admission transactions, 24,548 simulator-oracle comparisons, 1,000,000 allocation attempts, 1,000,000 physical DMA transfers, 700,000 stale-event injections, or 7,500 recovery and fallback episodes. All 800 omitted-cache-operation controls were detected. A 24-minute run completed 6,685,424 physical DMA lifecycle iterations without an observed data, device-call, or lifecycle failure. The timing audit also exposed a material contract boundary: measured CPU and RVV maxima of 19.926 and 19.778 μ s exceeded their registered 10 and 4 μ s bounds. Thus, AIRTOS supports evidence-governed admission and finite-domain runtime invariants on the tested configuration while correctly withholding a hard-real-time conclusion when the plan’s timing evidence is inapplicable.

Index Terms—Edge AI, heterogeneous computing, admission control, real-time systems, runtime assurance, DMA coherency, fault recovery.

I. INTRODUCTION

EDGE inference increasingly crosses a qualification boundary that conventional deployment interfaces leave implicit. A compiler emits an execution plan, yet the plan enters a live system whose provider health, contiguous memory, cache ownership, outstanding device commands, and admitted workload can differ from the conditions under which that plan was produced. Moreover, an inference is not necessarily one schedulable thread. It can be a directed acyclic graph (DAG) of CPU, vector, accelerator, and DMA segments with resource-specific execution bounds and asynchronous completions. Structural validity and numerical correctness alone therefore do not establish that a plan is eligible to run in the current state.

The broader toolchain has three responsibilities. Compilation constructs an execution plan for a target chip. Evidence production states the conditions under which that plan is

correct and the inputs to which it applies. AIRTOS provides the third stage: runtime governance between the qualified plan and physical execution. In this stage, a plan is not dispatched merely because an inference function exists. It must establish current eligibility with respect to its evidence, invocation, providers, memory, other deadline commitments, coherency actions, and recovery policy. AIRTOS therefore acts as a consumer-side admission boundary rather than another compiler pass or a new general-purpose kernel.

Real-time scheduling theory makes the assumptions behind timing guarantees explicit. Classical EDF results and sporadic-task feasibility analysis depend on bounded execution and defined arrival models [1], [2]; execution-time assurance work further shows that the strength of a timing claim is inseparable from the bounds supplied to the scheduler [3]. Heterogeneous task runtimes and schedulers already express DAG dependencies, locality, and resource types [4]–[8]. Accelerator-sharing systems likewise address preemption, spatial partitioning, memory pressure, and multi-model quality of service [9]–[12]. These mechanisms provide essential foundations, but they normally begin after the workload and its execution bounds have been accepted as valid inputs.

AIRTOS evaluates a joint predicate over parsing, binding, applicability, evidence, providers, memory, coherency, timing, and recovery. Admission probes an arena lease, captures the current schedule, evaluates a conservative multi-resource simulation, and commits memory and job state only if their generations remain unchanged. During execution, dependency-ready segments enter stable per-resource EDF queues. Plan-declared ownership transitions govern cache and DMA operations, while a device epoch and dispatch cookie scope completion events. Recovery consumes a bounded policy, quarantines an exhausted resource, and subjects a fallback to the same gate before execution.

The design follows the separation principle used in runtime assurance: a complex producer may propose an action, but a smaller consumer-side mechanism decides whether that action is admissible in the present state [15], [16]. Trace records are similarly restricted. They can select a subsequent compiler experiment, but they cannot elevate evidence or install a plan directly; every candidate returns through verification and admission. This preserves the distinction between observing an execution and establishing a reusable claim [17].

We evaluate AIRTOS through four claim-oriented cores and a short physical HIL run. The corpus covers qualification decisions, concurrent memory–schedule transactions, an independent scheduling oracle, arena leases, physical DMA with negative coherency controls, stale-event and recovery transitions, fallback gates, and trace classification. The timing

audit is equally important: CPU and RVV maxima exceeded their registered execution bounds. AIRTOS consequently identifies the tested plan as outside its timing contract rather than converting favorable percentiles into a hard-real-time claim.

This work makes four contributions:

- 1) an evidence-constrained admission predicate that combines plan qualification with live resource, memory, coherency, timing, and recovery state;
- 2) an optimistic atomic lease-and-schedule transaction that prevents half-admitted jobs under concurrent submissions;
- 3) dependency-aware resource governance with explicit ownership transitions and epoch-scoped completion and recovery semantics; and
- 4) a preregistered finite-domain evaluation that reports invariant checks, physical coherency evidence, and a negative timing-contract result.

The contribution is this governed composition and its consumer-checkable boundaries. EDF, heterogeneous DAG execution, accelerator sharing, DMA ownership, and timeout handling remain established mechanisms.

II. BACKGROUND AND RELATED WORK

A. Real-Time and Heterogeneous DAG Scheduling

Liu and Layland establish the classical uniprocessor foundation for EDF under periodic-task assumptions [1], while processor-demand analysis extends feasibility reasoning to sporadic workloads [2]. Vestal makes execution-time assurance an explicit scheduling dimension [3]. AIRTOS retains this discipline: its deadline statement is conditional on registered segment costs, non-preemptive blocking, modeled arrivals, coherency charges, and recovery charges. Its simulator rechecks all admitted deadlines after candidate insertion instead of treating EDF priority as a timing proof.

HEFT schedules precedence-constrained tasks across heterogeneous processors using rank-based list scheduling [4]. StarPU combines task scheduling with heterogeneous data management [5]; Legion expresses independence and locality through logical regions [6]; and PTask makes accelerator work and dataflow visible to the operating system [7]. Type-aware federated scheduling gives real-time semantics to typed heterogeneous DAGs [8]. AIRTOS does not present DAG queues or per-resource EDF as new algorithms. It places these mechanisms downstream of an evidence and applicability gate and couples their result to a generation-checked memory transaction and recovery state.

B. Accelerator Runtimes and Multi-Tenant AI

Accelerator sharing spans different hardware assumptions. PREMA exploits preemptible NPU execution [9], whereas Planaria partitions an accelerator spatially among tenants [10]. MoCA adapts execution around multi-tenant memory behavior [11], and DREAM targets dynamic real-time multi-model edge workloads [12]. Salus exposes fine-grained GPU sharing primitives and distinguishes persistent from transient memory [13]; Clockwork pursues predictable DNN serving through

controlled execution and explicit timing information [14]. These systems establish that scheduling, memory, and latency must be governed jointly. Their admitted unit, however, is generally a model or execution request rather than a plan carrying separately checkable domain and evidence obligations.

Embedded inference systems demonstrate how tightly runtime design is coupled to constrained memory and platform support. TensorFlow Lite Micro provides an embedded inference runtime for small devices [22], and MLPerf Tiny establishes reproducible measurement practices for such platforms [21]. AIRTOS focuses on the qualification layer: whether a particular evidence-carrying plan may enter the live runtime and how its memory, timing, ownership, and recovery obligations remain governed.

C. Runtime Assurance, Coherency, and Recovery

Simplex separates a high-performance controller from a trusted baseline and decision module [16]; modern runtime-assurance work generalizes that structure as a safety filter around complex components [15]. Runtime verification checks execution events against specified properties [17]. AIRTOS applies the separation at the compiler–runtime boundary: a producer supplies a plan and evidence, while the runtime checks current eligibility. Its trace can only propose an experiment, preventing an execution from certifying its replacement.

DMA specifications already define fences and ownership transitions for asynchronous devices [18], [19]. AIRTOS does not redefine these primitives. Instead, the plan names the buffer range and required clean, barrier, and invalidate actions; the runtime confines the range to the active lease and invokes provider hooks at the ownership boundary. Completion attribution adds an orthogonal device-generation rule: an event is accepted only when resource, epoch, active command, and cookie agree. The evaluation uses an independent scheduling oracle rather than deriving expected answers from the runtime under test, following the principle that oracle construction is part of the validity argument [20].

III. SYSTEM MODEL AND PROBLEM FORMULATION

A. Jobs, Plans, and Runtime State

An admitted job is

$$J_i = (id_i, P_i, r_i, d_i, \kappa_i, \chi_i), \quad (1)$$

where P_i is a compiler-produced plan, r_i is release time, d_i is an absolute deadline, κ_i is the minimum evidence policy, and χ_i is a recovery policy. A plan contains a segment DAG $G_i = (V_i, E_i)$. Each segment is

$$s = (id, res, C, Pred, off, size, flags), \quad (2)$$

where $res \in \{CPU, RVV, NPU, DMA\}$, C is a registered execution bound, $Pred$ is the predecessor set, off and $size$ identify a lease-relative buffer range, and $flags$ declare ownership actions. A segment becomes ready only after release and completion of every predecessor:

$$Ready(s, t) \Leftrightarrow t \geq r_i \wedge \forall u \in Pred(s), \text{ state}(u) = done. \quad (3)$$

The runtime state is

$$R_t = (\{Q_r\}, \{A_r\}, \mathcal{L}, \mathcal{C}, \{epoch_r\}, \{cookie_r\}, D_t, Z_t), \quad (4)$$

where Q_r and A_r are the ready queue and active segment for resource r , $\mathcal{L}$ is the active lease set, $\mathcal{C}$ is buffer ownership state, D_t records provider health and recovery, and Z_t is the attributable trace. Each queue is ordered by the owning job's absolute deadline with stable FIFO tie-breaking. Running device segments are non-preemptive.

B. Evidence-Constrained Admission

AIRTOS admits a candidate only when nine conditions hold jointly:

$$\begin{aligned} \text{Admit}(J_i, R_t) \Leftrightarrow & \text{Parse} \wedge \text{Bind} \wedge \text{Domain} \wedge \text{Evidence} \\ & \wedge \text{Provider} \wedge \text{Memory} \wedge \text{Coherence} \\ & \wedge \text{Sched} \wedge \text{Recoverable}. \end{aligned} \quad (5)$$

Parse covers safe package decoding. *Bind* links package, plan, model, target, runtime ABI, provider ABI, and fallback identities. *Domain* covers shape, dtype, layout, and byte ranges. *Evidence* requires every selected obligation to have the expected scope, artifact, verifier identity, resource, and verified status. *Provider* requires registered and healthy resources. *Memory* requires a non-overlapping lease that contains all declared ranges. *Coherence* requires supported ownership actions. *Sched* requires all old and new jobs to meet their deadlines in simulation. *Recoverable* requires valid budgets and a resource that is neither pending recovery nor quarantined.

The scheduling term is evaluated by SimEDF^+ , a discrete-event simulation with per-resource non-preemptive EDF, DAG releases, active residuals, coherency and recovery charges, and reservation demand. For a running segment started at t_s , the charged residual is

$$C_s^{\text{rem}} = \max(0, C_s - (t - t_s)) + O_s^{\text{coh}} + O_s^{\text{rec}}. \quad (6)$$

The candidate passes only if

$$\forall J_j \in \mathcal{J}_{\text{admitted}} \cup \{J_i\}, \quad \hat{F}_j \leq d_j, \quad (7)$$

and reservation demand at each checked horizon does not exceed supply. The all-job test is necessary because an earlier-deadline candidate can finish on time while delaying a previous commitment beyond its deadline.

C. Safety Properties and Conditional Timing

Lease isolation: if leases are pairwise disjoint, the loader confines segment ranges to the plan arena, and providers access only declared ranges, then segments from different sessions address disjoint intervals. This property concerns inter-session isolation; within-session alias validity remains part of the compiler plan.

DAG order: if a segment enters a resource queue only after all predecessors complete and only queued segments dispatch, every dispatch sequence is a topological extension of the plan DAG. EDF orders ready segments but does not bypass dependencies.

Epoch isolation: a completion event carries (r, e, c, status) . It is accepted only if

$$\text{Accept}(\text{irq}, R_t) \Leftrightarrow e = epoch_r \wedge c = cookie_r \wedge A_r \neq \emptyset. \quad (8)$$

A reset advances the epoch, so an event from an earlier generation cannot mutate a new command. Valid completion clears active association, rejecting a same-epoch duplicate; cookie wrap advances the epoch before reuse.

Conditional data consistency: for a non-coherent buffer, AIRTOS models

$$\begin{aligned} \text{CPU}_{\text{dirty}} & \xrightarrow{\text{clean+barrier}} \text{Device}, \\ \text{Device} & \xrightarrow{\text{complete+invalidate}} \text{CPU}_{\text{valid}}. \end{aligned} \quad (9)$$

If the plan declares the correct range and actions, provider hooks implement them, and the device obeys completion semantics, device reads observe submitted CPU data and later CPU reads observe device output.

Conditional deadline safety: let a segment dispatch at a_s , complete in hardware at h_s , and have registered budget C_s . AIRTOS exposes logical completion at

$$\ell_s = \max(h_s, a_s + C_s). \quad (10)$$

Holding early completion to the modeled boundary keeps successor releases and non-preemptive occupancy aligned with SimEDF^+ . If actual execution and charged overhead do not exceed registered bounds, arrivals and faults remain inside reservation, all jobs pass atomic admission, and runtime tie-breaking matches simulation, the schedules coincide by induction over events. A passing all-job check then preserves the deadline invariant. An execution exceeding its registered bound leaves this timing domain and requires requalification.

Atomic publication: consider a successful submission at the locked publication step. The lease generation and schedule generation have just been revalidated, the arena commit has succeeded, and no subsequent operation before job publication can fail. Hence a visible job owns a committed lease, while a rejected or retried submission publishes neither. The generation checks extend this argument to concurrent admissions by invalidating every stale probe or schedule snapshot before publication.

Bounded recovery closure: if cancellation, reset, and reinitialization polls respect their declared budgets and at most K_r reset attempts are permitted, a recovering resource reaches either healthy or quarantined state in finite model time. Retaining the failed job's lease until closure prevents the same range from being reassigned while an old device command can still exist.

IV. AIRTOS DESIGN

A. Evidence Flow and Governance

AIRTOS consumes an AEG v2 package containing plan and evidence identities, policy and target bindings, input domain, segment DAG, execution costs, arena size, ownership actions, reservation parameters, fallback identity, and recovery budget. Session creation validates structure before accessing payload objects, compares the package against a consumer-side

TABLE I
PLAN CONTRACT AND CONSUMER-SIDE ENFORCEMENT

Contract dimension	Runtime enforcement
Package and identities	Bounds-safe load; plan/model/target/ABI and fallback binding
Input and evidence	Shape/dtype/layout/bytes; obligation scope, artifact, verifier, resource, status
Resources and memory	Provider health; lease range, generation, and segment bounds
Timing and ownership	All-job $SimEDF^+$; declared clean/barrier/invalidate actions
Recovery and feedback	Timeout/attempt budget, quarantine, fallback re-admission, non-promoting trace




Fig. 1. AIRTOS places a consumer-side governance boundary between an evidence-carrying compiler plan and heterogeneous hardware. The joint gate qualifies the plan against current invocation and runtime state; attributed trace can propose a new experiment but cannot elevate evidence.

trust bundle, and checks each evidence obligation independently. Submission adds invocation-specific checks and current provider, memory, schedule, and recovery state.

Figure 1 makes the producer–consumer boundary explicit. No successful check substitutes for another: a healthy provider cannot compensate for invalid evidence, and a feasible schedule cannot compensate for an unavailable arena or unsupported ownership transition.

B. Atomic Lease and Schedule Admission

Memory and time are admitted as one optimistic transaction. Under the runtime lock, submission selects a free job slot, probes a first-fit interval, captures the lease generation, snapshots admitted work, and captures the schedule generation. It releases the lock while $SimEDF^+$ evaluates the candidate. On reacquiring the lock, it verifies that the slot and session remain available, provider health remains valid, and both generations match. Only then does it commit the lease, initialize the job, publish the slot, and advance schedule state.

If simulation rejects the candidate, no lease is committed. If either generation changes, submission retries from fresh state. If provider health changes at the final check, the committed lease is released before rejection. No fallible operation remains

between successful arena commit and job publication. This protects against both a concurrent allocation invalidating a probed interval and a concurrent admission or completion invalidating the schedule snapshot.

C. Resource and Coherency Governance

Each resource owns a stable EDF queue. A segment enters only when all predecessors are done. When a resource is idle, AIRTOS dispatches the head, assigns a cookie, records the epoch, performs pre-submit ownership actions, and invokes the provider. Independent segments on different resources may progress concurrently, while each active resource segment remains non-preemptive.

Arena leases define session-level physical isolation. Segment offsets are relative to the active lease, and coherency operations are rejected if their rounded cache-line range escapes it. Before submission AIRTOS optionally cleans the declared range and issues a barrier. After accepted completion it issues a barrier, optionally invalidates the range, and issues a final barrier. The plan records ownership intent; the provider implements the platform operation.

D. Epoch-Scoped Recovery and Feedback

A timeout first initiates cancellation. If cancellation does not confirm quiescence within budget, recovery advances the epoch and begins reset, followed by reinitialization and health checking. Attempts are bounded by the package policy. Exhaustion moves the resource to *Quarantined*, excluding it from provider admission. For cancellation budget Δ_c , reset budget Δ_r , reinitialization budget Δ_i , and at most K_r attempts, model-level closure is bounded by

$$T_{close} \leq \Delta_c + K_r(\Delta_r + \Delta_i) + O(K_r). \quad (11)$$

The failed job retains its lease while recovery remains open. Fallback is not an unconditional branch: AIRTOS re-evaluates trust and evidence, confirms the original lease is active and covers all fallback ranges, checks fallback provider health, snapshots current work, and invokes $SimEDF^+$ with the original absolute deadline. Only a passing fallback returns to pending execution.

Trace entries include logical sequence, timestamp, run identifier, plan, job, segment, resource, epoch, cookie, event, status, and queue depth. A classifier maps attributed symptoms to candidate experiments. As summarized in Figure 2, trace cannot modify a plan or evidence status. Every candidate receives a new identity and returns through compiler verification and AIRTOS admission, avoiding circular self-certification.

V. IMPLEMENTATION

A. Production Runtime Paths

The implementation is a portable C governance layer integrated with RT-Smart rather than a replacement kernel. Production modules separate package parsing, evidence evaluation, session/submission management, schedule simulation, coordination, memory leasing, coherency, recovery, and tracing.




Fig. 2. Atomic admission and recovery flow. Lease and schedule generations guard publication; resource, epoch, and cookie guard completion; bounded recovery ends in health or quarantine, and fallback re-enters the full gate.

The loader validates section bounds, counts, dependencies, resources, hashes, numeric ranges, recovery policy, and primary/fallback structure before building a session. Evidence evaluation compares plan, policy, model, target, runtime-ABI, and provider-ABI identities and iterates obligation records, including artifact and verifier identities and allowlisting. Submission performs domain, interarrival, deadline, provider, lease, and schedule checks. The simulator reconstructs admitted jobs and the candidate in independent state, including active residuals and reservations.

Runtime queues maintain dependency-ready order. The coordinator is invoked through a polling entry point and serializes state transitions with the platform lock. Completion entry points apply Equation (8) before post-device ownership actions and successor release. The allocator uses generation-checked first fit over a fixed arena. Recovery records per-resource cancel, reset, reinitialization, health, attempt count, and quarantine states. The trace ring uses monotonically increasing logical sequence numbers and event-level primary or fallback plan identity.

B. K230 Integration and Tested Plan

The physical evaluation uses one CanMV-K230-LP4 V3.0 board running RT-Smart. The fixed plan performs float32 Add+ReLU on shape [1, 8] in a 64-byte arena, with a 100 μ s relative deadline and 100 μ s minimum interarrival. The primary RVV segment registers a 4 μ s execution bound, 1 μ s coherency charge, and 50 μ s recovery charge. CPU fallback registers 10, 1, and 50 μ s, respectively.

Physical ownership tests allocate cached MMZ buffers, obtain physical addresses, execute writeback and invalidation through the platform interface, and move data through GSDMA. The complete path generates a CPU pattern, writes back the source, performs physical copy, invalidates the destination, and compares bytes. Device lifecycle testing deinitializes and reinitializes GSDMA and verifies a subsequent copy. The resource model can express an NPU when an appropriate

provider and evidence exist; the reported physical experiment uses the fixed CPU/RVV/DMA workload and does not report general NPU performance.

VI. EXPERIMENTAL METHODOLOGY

A. Preregistered Questions and Oracles

The evaluation is organized by claims. Core 1 asks whether qualification and admission fail closed across package, binding, input, evidence, provider, memory, coherency, timing, recovery, health races, trust rotation, and concurrent admission. Unsafe admission, overlap, partial commit, or rollback residue is a zero-tolerance endpoint.

Core 2 asks whether scheduling decisions match an independent oracle and whether the registered timing contract applies on the board. The oracle does not call production scheduler, queue, or lease code; it compares status and finish time for frozen DAG scenarios. Software-model admission ablations and execution-time sensitivity are reported separately from board timing.

Core 3 asks whether lease and ownership invariants hold. Allocation is checked against independent ownership state; physical output is compared byte for byte. Omitted-clean and omitted-invalidate controls must be observable. Core 4 asks whether old-world events are isolated and recovery terminates at a governed state. Stale classes, bounded recovery, quarantine, fallback re-admission, trace classification, physical lifecycle, and short continuous operation are evaluated separately. Classification thresholds are macro-F1 ≥ 0.90 , top-3 recall ≥ 0.95 , and gate bypass = 0.

B. Platform and Measurement Discipline

Protocol `airtos-exp-v6` was executed on 2026-08-05. The loader and scheduling corpus is 49,085,520 bytes. Physical transfer sizes are 64, 256, 4,096, and 65,536 bytes, each receiving 250,000 transfers. The continuous run lasts 1,440 s, cycles the same sizes, and reinitializes the device every 100,000 iterations.

Schedule agreement concerns frozen simulator scenarios, not physical execution of 24,548 heterogeneous workloads. Stale and recovery cases execute production state-machine code on the board through controlled in-process callbacks; they are distinct from naturally late driver interrupts. Device reopen/reinitialize calls the physical library and is distinct from chip hard reset. Timing is audited against bounds already registered in the plan. The recomputed summary and its SHA-256 digest are retained in the reproducibility artifact.

VII. EVALUATION

A. Admission and Scheduling Consistency

Table III summarizes Core 1. The admission matrix completed 3,900 cases without an observed decision failure. Diagnostic classification completed 23,400 cases with macro-F1 1.0. Provider-health races and trust rotation added 300 and 1,500 decisions without unsafe commit or decision failure. Across 400,000 concurrent transactions, 112,091 were accepted and 287,909 rejected. The ratio describes a validation corpus

TABLE II
EVIDENCE LAYERS AND THE STATEMENTS THEY SUPPORT

Layer	Executed object	Supported statement
Software-model corpus	Production loader, admission transaction, <i>SimEDF</i> ⁺ , allocator, recovery, and trace logic	Finite-domain implementation consistency, all-job admission ablation, WCET sensitivity, transaction atomicity, and state-machine decisions
Physical K230 data path	Cached MMZ buffers, cache operations, physical addresses, and GSDMA	Data consistency for the tested engine, aligned buffers, and four transfer sizes, with observable clean/invalidate omission controls
Controlled board injection	Production completion/recovery paths with in-process stale and fault callbacks	Epoch-cookie rejection, budget/quarantine closure, and fallback-gate behavior for the injected event classes
Short continuous HIL	Repeated physical GSDMA lifecycle with deinitialize/reinitialize and byte comparison	No observed data, device-call, or lifecycle counterexample in 6,685,424 iterations over 1,440 s




Fig. 3. Finite-domain operation coverage across the four experimental cores. Counts aggregate non-overlapping primary operations for visualization; they are not samples from a common estimator. Each core recorded zero observed safety failures under its defined endpoint.

TABLE III
EVIDENCE-CONSTRAINED AND ATOMIC ADMISSION

Endpoint	Cases	Failures
Admission matrix	3,900	0
Diagnostic classification	23,400	0
Provider-health race	300	0
Trust-root rotation	1,500	0
Concurrent transactions	400,000	0
Active lease overlap	400,000	0
Partial commit	400,000	0

rather than workload capacity. No active overlap or partial commit was observed.

Core 2 executed 7,950 loader cases and 24,548 oracle comparisons. Loader decisions matched frozen expectations, and simulator status and finish time matched the oracle in every scenario. The software-model ablation in Table IV shows why admission must include all jobs. With no check, 5,822 of 10,000 scenarios were falsely accepted. Checking only candidate finish still produced 3,932 false accepts because a candidate could delay an existing commitment. FIFO and fixed-priority admission reduced false accepts but added false rejects. Complete all-job *SimEDF*⁺ produced zero false accept and zero false reject against the oracle in this corpus.

TABLE IV
SOFTWARE-MODEL ADMISSION ABLATION

Condition	False accept	False reject
No admission	5,822	0
Candidate only	3,932	0
FIFO	27	355
Fixed priority	29	380
All-job <i>SimEDF</i> ⁺	0	0

A sensitivity experiment held the set at 4,178 admitted scenarios per ratio. At actual/WCET ratios 0.50, 0.80, and 1.00, no deadline miss was observed. At 1.05 and 1.20, misses rose to 446 and 1,145. This controlled transition supports the registered bound as a premise of admission rather than a descriptive performance target. CPU and RVV Add+ReLU each completed 30,000 calls with zero numerical failure for the fixed operator and shape.

B. Memory Isolation and Physical Coherency

The allocator completed 1,000,000 attempts, of which 948,950 produced leases. Independent checks observed no overlap, corruption, generation error, or rollback failure counted by the frozen endpoint. Allocation failure under pressure is an availability outcome, not a safety failure.

The K230 path completed 1,000,000 physical DMA transfers with zero byte mismatch. Each size contributed 250,000 transfers. Mean complete-path time was 29.949, 34.490, 86.502, and 713.355 μ s for 64, 256, 4,096, and 65,536 bytes, respectively. Omitting source clean was detected in 400/400 controls, and omitting destination invalidate was detected in 400/400. Thus, the complete path is tested on an environment where both ownership operations are observable. The conclusion applies to the aligned buffers, four sizes, cache interface, and GSDMA engine used on this K230 configuration.

C. Recovery, Feedback, and Short HIL

Seven stale-event classes contributed 700,000 controlled events, with no observed invalid state mutation. Recovery completed 1,500 base episodes, and 4,800 budget episodes checked attempt accounting, closure, and quarantine. Another 1,200 fallback episodes exercised trust, evidence, lease, provider, and schedule gates without a re-admission bypass.

The trace classifier completed 800 base cases with macro-F1 and top-3 recall both 1.0. A robust corpus added 2,400 noise




Fig. 4. K230 timing-contract audit. Although batch p99 is below each registered bound, the observed maximum exceeds both CPU and RVV WCETs; maximum-based timing applicability therefore fails for both paths.

and ring-wrap cases and retained macro-F1 1.0. These values establish behavior on the eight frozen synthetic classes and confirm stable decisions under irrelevant events and ring wrap.

Physical GSDMA reopen/reinitialize completed 300 episodes with zero lifecycle failure; p99 was 53.148 μ s and maximum 69.222 μ s. The short HIL run completed 6,685,424 DMA lifecycle iterations in 1,440 s. Each iteration generated a pattern, selected one of four sizes, wrote back the source, invoked GSDMA, invalidated the destination, and compared bytes. Reinitialization occurred every 100,000 iterations. The run recorded zero data, device-call, and lifecycle failures. Temperature ranged from 49.103 to 52.706 $^{\circ}$ C, ending at 52.406 $^{\circ}$ C. These are physical DMA lifecycle iterations rather than NPU executions or heterogeneous inference jobs.

The operation totals in Figure 3 are 429,100 for Core 1, 92,498 for Core 2, 2,000,800 for Core 3, and 7,396,424 for Core 4. They visualize coverage but are not combined into a common failure-rate estimate.

D. Timing Contract Audit

Figure 4 and Table V separate percentile behavior from contract validity. Both execution paths had maximum batch p99 of 1.592 μ s. The RVV maximum was 19.778 μ s against a registered 4 μ s bound, and the CPU maximum was 19.926 μ s against 10 μ s. Both contracts therefore fail on maximum observation. This is the intended result of evidence-constrained governance: a plan may remain structurally valid and numerically correct while its timing evidence is rejected for the measured platform state.

The maximum batch p99 among seven control operations was 3.741 μ s, so the preregistered deadline-relative 5 μ s criterion passed. The stricter 0.2 μ s criterion, equal to 5% of the shortest 4 μ s segment, did not pass. AIRTOS fits the deadline-relative control budget used by the protocol, but the current implementation/plan pair does not support negligible overhead relative to its shortest segment.

TABLE V
K230 TIMING-CONTRACT AUDIT

Path	Bound	p99 (μ s)	Maximum
RVV Add+ReLU	4	1.592	19.778
CPU Add+ReLU	10	1.592	19.926
Clock read	5	1.518	12.111
ISR completion	5	3.445	39.074
Lease release	5	1.667	2.629
Plan load	5	3.741	36.963
Queue push/pop	5	1.555	37.037
<i>SimEDF</i> ⁺ admission	5	2.223	37.037
Trace emission	5	1.704	35.777

The timing result does not contradict the conditional scheduling argument. The oracle comparisons establish that the implementation evaluates its declared model consistently. Board measurement establishes that the declared 4 and 10 μ s bounds are not valid for the observed physical domain. Recalibration and a new evidence-carrying plan are consequently required before a board-level deadline claim can be evaluated.

VIII. DISCUSSION

The combined results support three conclusions about the tested artifact. First, qualification is operational rather than documentary: package, evidence, domain, provider, memory, schedule, coherency, and recovery checks all participate before dispatch. The concurrent corpus supports the intended linearization of memory and schedule state within the tested interleavings.

Second, ownership is explicit across both memory and device generations. Arena ranges attach to sessions, cache transitions attach to plan segments, and completions attach to epoch and cookie. The physical DMA result and omission controls connect abstract ownership to observable data on the tested K230 configuration. The controlled stale-event corpus connects generation matching to the production recovery state machine.

Third, evidence remains capable of rejecting an otherwise functioning plan. The Add+ReLU paths were numerically correct in 60,000 calls and their batch p99 values were favorable, yet maxima exceeded registered bounds. Withholding a deadline claim demonstrates the distinction between functional success and timing applicability that motivates Equation (5). Trace preserves the same discipline: classification can select an experiment, while a new plan identity, new evidence, and fresh admission prevent online self-certification.

The physical scope is one CanMV-K230-LP4 V3.0 board, RT-Smart, a fixed float32 Add+ReLU plan, CPU/RVV execution, and GSDMA over four aligned sizes. The physical coherency statement applies to that buffer, cache-interface, engine, and size contract. The corpus does not provide general NPU performance, board-to-board variability, or power data. Controlled stale events exercise production state-machine code but are distinct from naturally delayed driver interrupts; reopen/reinitialize is a physical library lifecycle operation, not chip hard reset. These distinctions maintain correspondence between mechanism, stimulus, and supported statement.

The continuous run provides 24 minutes and 6,685,424 verified DMA lifecycle iterations, during which no recorded data, device-call, or lifecycle counterexample occurred. It is not used as a substitute for a 24-hour study. The deadline theorem remains conditional on registered execution, arrival, blocking, coherency, and recovery bounds. The negative WCET audit identifies which premise does not apply to the current plan; a recalibrated plan can pass through the same consumer-side machinery without changing admission semantics.

IX. CONCLUSION

AIRTOS treats an evidence-carrying heterogeneous execution plan as a first-class runtime object. Its joint predicate binds compiler evidence and applicability to live provider, memory, coherency, schedule, and recovery state. Generation-checked lease and schedule commit prevents half-admission, dependency-ready queues govern execution, plan-declared ownership actions connect buffers to physical DMA, and epoch-cookie matching separates current commands from stale completions. Bounded recovery ends in restored health or quarantine, while fallback and trace-generated candidates return through the complete gate.

Across the frozen corpus and one K230 board, these mechanisms produced no observed safety failure at the registered endpoints, including one million physical transfers and 6.7 million short-run lifecycle iterations. At the same time, measured maxima invalidated both registered execution bounds. The result illustrates the role of evidence-constrained governance: the runtime does not infer eligibility from structural validity, numerical success, or favorable percentiles. It admits, rejects, recovers, and traces plans according to explicit, consumer-checkable conditions.

REFERENCES

- [1] C. L. Liu and J. W. Layland, "Scheduling algorithms for multiprogramming in a hard-real-time environment," *J. ACM*, vol. 20, no. 1, pp. 46–61, 1973.
- [2] S. Baruah, A. K. Mok, and L. E. Rosier, "Preemptively scheduling hard-real-time sporadic tasks on one processor," in *Proc. IEEE Real-Time Syst. Symp.*, 1990, pp. 182–190.
- [3] S. Vestal, "Preemptive scheduling of multi-criticality systems with varying degrees of execution time assurance," in *Proc. IEEE Real-Time Syst. Symp.*, 2007, pp. 239–243.
- [4] H. Topcuoglu, S. Hariri, and M.-Y. Wu, "Performance-effective and low-complexity task scheduling for heterogeneous computing," *IEEE Trans. Parallel Distrib. Syst.*, vol. 13, no. 3, pp. 260–274, 2002.
- [5] C. Augonnet, S. Thibault, R. Namyst, and P.-A. Wacrenier, "StarPU: A unified platform for task scheduling on heterogeneous multicore architectures," *Concurrency Computat.: Pract. Exper.*, vol. 23, no. 2, pp. 187–198, 2011.
- [6] M. Bauer, S. Treichler, E. Slaughter, and A. Aiken, "Legion: Expressing locality and independence with logical regions," in *Proc. Int. Conf. High Perform. Comput., Netw., Storage Anal.*, 2012, pp. 1–11.
- [7] C. J. Rossbach, J. Currey, M. Silberstein, B. Ray, and E. Witchel, "PTask: Operating system abstractions to manage GPUs as compute devices," in *Proc. 23rd ACM Symp. Operating Syst. Princ.*, 2011, pp. 233–248.
- [8] C.-C. Lin, J. Shi, N. Ueter, M. Guenzel, J. Reineke, and J.-J. Chen, "Type-aware federated scheduling for typed DAG tasks on heterogeneous multicore platforms," *IEEE Trans. Comput.*, vol. 71, no. 12, pp. 3266–3280, 2022.
- [9] Y. Choi and M. Rhu, "PREMA: A predictive multi-task scheduling algorithm for preemptible neural processing units," in *Proc. IEEE Int. Symp. High Perform. Comput. Archit.*, 2020, pp. 220–233.
- [10] S. Ghodrati *et al.*, "Planaria: Dynamic architecture fission for spatial multi-tenant acceleration of deep neural networks," in *Proc. 53rd IEEE/ACM Int. Symp. Microarchitecture*, 2020, pp. 681–697.
- [11] S. Kim *et al.*, "MoCA: Memory-centric, adaptive execution for multi-tenant deep neural networks," in *Proc. IEEE Int. Symp. High Perform. Comput. Archit.*, 2023, pp. 828–841.
- [12] S. Kim, H. Kwon, J. Song, *et al.*, "DREAM: A dynamic scheduler for dynamic real-time multi-model ML workloads," in *Proc. Int. Conf. Hardware/Software Codesign Syst. Synth.*, 2023, pp. 1–12.
- [13] P. Yu and M. Chowdhury, "Salus: Fine-grained GPU sharing primitives for deep learning applications," in *Proc. Mach. Learn. Syst.*, 2019, pp. 98–111.
- [14] A. Gujarati, R. Karimi, S. Alzayat, *et al.*, "Serving DNNs like Clockwork: Performance predictability from the bottom up," in *Proc. 14th USENIX Symp. Operating Syst. Design Implement.*, 2020, pp. 443–462.
- [15] K. L. Hobbs, M. Mote, M. Abate, S. Coogan, and E. Feron, "Runtime assurance for safety-critical systems: An introduction to safety filtering approaches for complex control systems," *IEEE Control Syst. Mag.*, vol. 43, no. 2, pp. 28–65, 2023.
- [16] D. Seto, B. H. Krogh, L. Sha, and A. Chutinan, "The Simplex architecture for safe online control system upgrades," in *Proc. Amer. Control Conf.*, 1998, pp. 3504–3508.
- [17] M. Leucker and C. Schallhart, "A brief account of runtime verification," *J. Logic Algebr. Program.*, vol. 78, no. 5, pp. 293–303, 2009.
- [18] Linux Kernel Documentation, "Buffer sharing and synchronization: DMA fences," 2026. [Online]. Available: <https://docs.kernel.org/driver-api/dma-buf.html>
- [19] Linux Kernel Documentation, "DMA API howto," 2026. [Online]. Available: <https://docs.kernel.org/core-api/dma-api-howto.html>
- [20] E. T. Barr, M. Harman, P. McMinn, M. Shahbaz, and S. Yoo, "The oracle problem in software testing: A survey," *IEEE Trans. Softw. Eng.*, vol. 41, no. 5, pp. 507–525, 2015.
- [21] C. R. Banbury, V. J. Reddi, P. Torelli, *et al.*, "MLPerf Tiny benchmark," in *Proc. NeurIPS Datasets Benchmarks Track*, 2021.
- [22] R. David, J. Duke, A. Jain, *et al.*, "TensorFlow Lite Micro: Embedded machine learning on TinyML systems," in *Proc. Mach. Learn. Syst.*, 2021, pp. 800–811.